

Advanced Programming

Testing

Fritz Henglein

Datalogisk Institut, Københavns Universitet (DIKU)

- Abstract data types and modules
 - Specifications
 - Testing fundamentals
 - Specification-driven testing
 - Property-based testing by input partitioning
 - Randomized testing
-

Goals

Week 5 (last week):

- concurrency;
- channels and threads;
- thread-based asynchronous programming;
- client-server programming.

Week 6 (this week):

- specification-driven testing;
- randomized property based testing;
- programming with Quickcheck.

Abstract data types and modules revisited

- An *abstract data type (ADT)* is an abstract type together with operations involving the type and a precise description of what the operations do. An ADT consists of
 - an *interface* containing the *declarations* of the type and its operations (functions),
 - their *specification* (what they do, which properties they have); and
 - *examples* illustrating correct and incorrect use of the operations.
- An *implementation* of an ADT contains
 - the *definitions* (source code) of the operations; and
 - *evidence* that it satisfies the ADT's specification (test design documentation and resulting test suite).
- An ADT typically has *multiple possible implementations*.
 - In particular, the implementation types may be different from each other.
 - Example: $s \rightarrow (a, s)$ and $\text{IORef } s \rightarrow \text{IO } a$ for state monads.
- *Module*: Structure with zero, one or more abstract or concrete type declarations.
 - ADT: Module with one abstract type declaration.

Abstract data types

- Programming language support for ADTs and *modular abstraction*:
 - ML-family languages: Signatures (module types), structures (module implementations), functors (module-parameterized structures).
 - Haskell: Type classes, class instances, class constraints in functions.
 - Java-like languages: Interfaces and abstract classes, concrete classes, concrete classes with interface-parameterized constructor(s).

Specifications

- *Specification*: Describes what a (collection of) function(s) does, not how it is implemented (its code).
 - Litmus test: Contains *everything* a user of a module *needs to know* and *nothing else*.
 - Specifications may be
 - inconsistent (contradictory): contain a contradiction;
 - incomplete (often the case): leave room for multiple implementations with different behaviors.
- *Implementation*: Code.
- *Correctness*: Evidence (e.g. test design method and resulting *test suite*) that code may satisfy its specification.
- *Requirements*: Informal description of desired functionality of a program component (what it should do and what it should not do).
 - Requirements may be unclear and leave room for interpretation.
 - A specification *resolves and documents* this ambiguity.

Module interface specification as a contract

- Module interface specification: *Contract* between programmer *using* a module (user) and programmer *implementing* it (implementor).
 - User to implementor: “All I want to know needs to be written in the module specification. How I use your implementation is none of your business. Don’t make me look at your code.”
 - Implementor to user: “All you need to know is written in the module specification. How I implement it is none of your business. Don’t ask to look at my code. (I may change it at any time.)”
 - Implementor and user are *roles*: They *may* be the same person.
- What if specification and implementation are *inconsistent*, as evidenced by a *test failure*?
 - Change the code to meet the specification: Requires no negotiation with user(s).
 - Change the specification to meet the code: Requires negotiation with user(s) since *changes may propagate to whatever code they have written*.
 - Change both, code and specification: Natural part of *exploratory programming* during *requirements elicitation*.
 - Common problem: Most modules are *underspecified*. User tries out code and relies on unspecified behavior of implementation.

Example: State monad specification

■ How to (partially) specify an abstract data type? Typical methods:

1 Properties its functions must satisfy.

■ Usually *universal equational properties*

2 Reference implementation (executable specification, model implementation)

■ Observable results must be the same as those from a reference implementation

■ Example: State monad

1 Monad laws plus

1	<code>put s >> get</code>	<code>= put s >> return s</code>
2	<code>put s >> put t</code>	<code>= put t</code>
3	<code>get >> get</code>	<code>= get</code>

for all s, t .

2 Coherence with purely functional implementation FState

Example: Reference implementation by type conversion

- Idea: Provide conversion functions from/to reference type
- Example:

```
1 FState s a = s -> (a, s)           -- reference type
2 IState s a = IORef s -> IO a       -- implementation type
3
4 to :: FState s a -> IState s a
5 to sa ref = do
6   s <- readIORef ref
7   let (a, t) = sa s
8   writeIORef ref t
9   return a
10
11 from :: IState s a -> (s -> IO (a, s))
12 from isa s = do
13   ref <- newIORef s
14   a <- isa ref
15   t <- readIORef ref
16   return (a, t)
```


Example: Reference implementation by type conversion

- Idea: Check that implementation functions fI simulate reference functions fF .
- Properties:

```

1  from (to sa)           ~ sa           :: FState s a
2  from getI             ~ getF          :: FState s s
3  from (putI s)          ~ putF s       :: FState s ()
4  from (returnI a)       ~ returnF a    :: FState s a
5  from (to sa 'bindI' (to . f) ~ sa 'bindF' f :: FState s b

```

for all $fstate$, s , a , f .

- Note: Left-hand side operations from new state monad implementation; right-hand side from reference implementation.
- Problems:
 - What does \sim mean?
 - How to generate functions for testing?

Testing

- *Testing* is (partial) verification by
 - executing code on inputs and checking whether the result is evidence of a specification violation;
 - *systematically* constructing the inputs to *maximize their likelihood of demonstrating a specification violation*.
- Terminology:
 - **Formal verification:** Mathematical proof that there exists *no* valid input that results in a specification violation.
 - **Testing:** The *systematic* discipline of finding “nasty” inputs that maximize the likelihood of finding specification violations.
 - Not being able to “break the code” is taken as evidence (not proof) that the code may be correct with respect to its given specification.
 - **Trying out, illustrating, exemplifying:** Applications of the code showing that it satisfies its specification for *some* inputs (if there is a specification) or that it does something useful (if there is no specification, only informal requirements).
- Don't use the term “test” unless you have *both* specification and code and you are *systematically* looking for specification violations.
 - Use “illustrate” or “exemplify” or “try out” instead.

Testing

- *Unit testing*: Testing a module implementation against its module specification.
 - A *unit* is a basic testable component; it corresponds to the notion of *module* here.
 - Testing is of the implementation including code in standard modules, but excluding bespoke modules it depends on. (These need to be stubbed in unit testing.)
 - NB: These notes are only about unit testing.
- *Integration testing*: Testing a composition of modules that form a *system component* against its specification.
- *System testing*: Testing a *deployable system* against its *specification* (which is typically just a large and changing test suite).
- *Acceptance testing*: Testing a deployable system against its *business and user requirements*.

Specification-driven testing

- Recall: All forms of testing require *both* specification and code.
 - No specification, no testing.
- *Specification-driven testing*: Analyze specification and systematically
 - find *valid* data with maximal likelihood to expose a specification violation of the code;
 - determine corresponding *expected* outputs *from the specification*.
 - Don't look at the code in this process.
- Test-driven development (TDD):
 - Develop the test suite *before* you have any code.
 - Continuously run the test suite to drive and check the code as you develop it.
- Specification-driven testing is also called *specification-based testing*, *external testing*, and *black-box testing*.
 - “Black-box testing” is not used here since much of questionable value is unfortunately written on the web on that.

Properties

- A *logical statement* is a statement that is either true or false.
- A *(logical) property* of one or more functions is a true logical statement involving the function(s).
- A *universally quantified property* has the form

$$\forall x, y, \dots Q(x, y, \dots)$$

or

$$\forall x, y, \dots (P(x, y, \dots) \Rightarrow Q(x, y, \dots)).$$

In words: For all (valid) x, y, \dots the property $Q(x, y, \dots)$ is true.

- Example:

$$\forall x \in \text{Int} ((x \geq 0) \Rightarrow \text{fib}(x) = fib(x))$$

where `fib` is some code and *fib* is the “true” mathematically defined Fibonacci function.

- A *(partial) formal specification* is often a conjunction of universally quantified properties.
- *Property-based testing*: Systematic design of *finite* subset of inputs aimed at falsifying a property.

Input partitioning

- *Input partitioning* is a design method for constructing *test data* by inspecting the *properties to be checked*.
- Given a universally quantified property, partition its valid inputs into a *finite* set of partitions, that is into pairwise disjoint subsets whose union is the entire set of valid inputs.
- From each partition,
 - choose a *typical* element and one or more *boundary elements*;
 - determine the *expected* output for each element according to the function specification—*not* the code.
- Test suite: The chosen input/expected output pairs. The boundary elements increase the likelihood of finding a specification violation.
- Test execution: Run the function on the inputs in the test suite; check that it produces the corresponding expected outputs.
- Input partitioning is also called *equivalence partitioning* or *equivalence class partitioning*.

Input partitioning: Numbers

- Generally useful partitioning of integral and floating-point numbers.
 - `int` (32-bit 2's-complement integers):
 - Partitions: Negative numbers, `{0}`, positive numbers.
 - Values: `-2147483648`, `-137`, `-1`, `0`, `1`, `2377`, `2147483647`.
 - `float` (64-bit IEEE 754 floating-point numbers):
 - Partitions: `{ Negative infinity }`, negative numbers, `{0}`, positive numbers, `{ positive infinity }`, `{ NaN }` ("not-a-number").
 - Values: `System.Double.NegativeInfinity`, `2.2250738585072014e-308`, `-1.0`; `-System.Double.Epsilon`, `0.0`, `System.Double.Epsilon`, `1.0`, `47771354343.98989988`, `1.7976931348623158e+308`, `System.Double.PositiveInfinity`, `System.Double.NaN`.
 - Note: F# syntax used!
- Note that most test values are boundary values.

Intermediate summary:

Programming = Specification + coding + verification

- *Programming* consists of

- 1 *Specification*: What is computed? Which functions, what are their *properties*?
- 2 *Coding*: How is it computed? How are the functions implemented (“code”)?
- 3 *Verification*: Which explicit *evidence* is there for the program code meeting its specification (correctness)?

- Verification always requires *both*, program code *and* specification.

- Program code may be correct with respect to one specification and incorrect with respect to another specification.

- Avoid statements such as “My code works”.

- With respect to which specification?

- Code is trivially correct with respect to whatever it happens to do—there is nothing to test; it does what it does. But what *does* it do that a code user can rely on?

- Does it *always* work and can you provide conclusive evidence for that, for *all* valid inputs and in *all* conceivable contexts, or just on *some data* that you happen to have tried out?

Property-based randomized testing

- Consider an implementation of a module specification.
- Let the module specification contain a *partial formal specification* in the form of a *conjunction of universal properties* that involve the module's functions:

$$\forall x, y, \dots. P_1(x, y, \dots) \wedge \dots \wedge \forall x, y, \dots. P_n(x, y, \dots)$$

where the P_i are *decidable properties*, that is there are terminating functions f_i returning a Boolean such that $f(x, y, \dots) == \text{True}$ if and only $P(x, y, \dots)$ holds.

- The module specification may contain additional requirements.
- *Randomized testing*: For each P_i
 - generate tuples $(x_0, y_0, \dots), \dots, (x_n, y_n, \dots)$ (*pseudo*)randomly.
 - apply $f_i(x_j, y_j, \dots)$ to all tuples (x_j, y_j, \dots) ;
 - If one application *fails* (returns *False*), this *proves* that the specification is *not satisfied* by the module implementation;
 - if all applications *succeed* (return *True*), this *provides empirical evidence* that the specification *may be satisfied* by the module implementation.

QuickCheck: Defining properties

```
quickCheck :: Testable a => a -> IO ()

class Testable a where
  property :: a -> Property

data Property

forall :: Testable b => Gen a -> (a -> b) -> Property
(==>) :: Testable a => Bool -> a -> Property
(==) :: Eq a => a -> a -> Property
(=/=) :: Eq a => a -> a -> Property
ioProperty :: Testable a => IO a -> Property
```

■ QuickCheck terminology:

- *Property*: Probabilistic computation that succeeds (returns true) or fails (returns false) with information about the failure.
- *Testable type*: Type with a default function for turning a value into a property.

QuickCheck: Probabilistic generation of values

```
newtype Gen a = MkGen {unGen :: QCGen -> Int -> a}
```

```
class Arbitrary a where  
  arbitrary :: Gen a
```

```
choose :: (a, a) -> Gen a  
oneof  :: [Gen a] -> Gen a  
frequency :: [(Int, Gen a)] -> Gen a  
elements :: [a] -> Gen a  
sized :: (Int -> Gen a) -> Gen a  
getSize :: Gen Int  
resize :: Int -> Gen a -> Gen a  
generate :: Gen a -> IO a
```

■ QuickCheck terminology:

- Gen a: Computations generating a (pseudo)random value of type a.
- Arbitrary: Class of types with a default random value generator for that type.

Input partitioning versus random testing

- Input partitioning: Good at identifying rare boundary elements/corner cases with high falsification potential.
- Random testing: Good at automatically generating large numbers of typical (random) elements.
- QuickCheck can be used to combine both: Custom generators (reflecting outcome of input partitioning)
- Explicit properties are crucial in both!

QuickCheck: More information

- Library randomized testing Haskell:
`https://hackage-content.haskell.org/package/QuickCheck-2.16.0.0/docs/Test-QuickCheck.html`
- QuickCheck manual: `https://www.cse.chalmers.se/~rjmh/QuickCheck/manual.html`

Testing challenges

- Testing imperative code (including object-oriented code):
 - Test design: Partition *state space*. Design *commands* for reaching boundary and typical states. Use *queries* for observing the reached state. Do the observed values have the expected *properties*?
 - Command-Query Separation (CQS) design pattern useful for testability.
- Testing concurrent code:
 - Properties need to be checked *for all possible schedulers* (all interleavings of all concurrently executing threads), not only for the particular scheduler one happens use.
 - Very difficult *and* expensive!
 - Advanced techniques (partially symbolic execution, etc) are beyond the scope of AP.
- Confusing test design with test suite
 - Specification-driven testing is the systematic and documented *process* of analyzing a specification to arrive at a *test suite with high falsification power*.
 - The test suite is the *output of the process*. By itself it says nothing about its falsification power. Only the *test design report* documenting the *systematic process* of construction the test suite does.

Summary

- Abstract data types and modules
- Specifications
- Testing fundamentals
- Specification-driven testing
- Property-based testing by input partitioning
- Randomized testing